

islr-ch4

Rohit Goswami

```
+++ title = "ISLR :: Classification" author = ["Rohit Goswami"] date = 2020-02-17T15:28:00+00:00 tags = ["solutions", "R", "ISLR"] categories = ["programming"] draft = false toc = true comments = true +++
```

Chapter IV - Classification

All the questions are as per the ISL seventh printing of the First edition [^fn:1].

Common Stuff

Here I'll load things I will be using throughout, mostly libraries.

```
libsUsed<-c("dplyr","ggplot2","tidyverse","ISLR","caret")
invisible(lapply(libsUsed, library, character.only = TRUE))
```

Question 4.10 - Page 171

This question should be answered using the `Weekly` data set, which is part of the ISLR package. This data is similar in nature to the `Smarket` data from this chapter's lab, except that it contains 1,089 weekly returns for 21 years, from the beginning of 1990 to the end of 2010.

- (a)** Produce some numerical and graphical summaries of the `Weekly` data. Do there appear to be any patterns?
- (b)** Use the full data set to perform a logistic regression with `Direction` as the response and the five lag variables plus `Volume` as predictors. Use the summary function to print the results. Do any of the predictors appear to be statistically significant? If so, which ones?
- (c)** Compute the confusion matrix and overall fraction of correct predictions. Explain what the confusion matrix is telling you about the types of mistakes made by logistic regression.
- (d)** Now fit the logistic regression model using a training data period from 1990 to 2008, with `Lag2` as the only predictor. Compute the confusion matrix and the overall fraction of correct predictions for the held out data (that is, the data from 2009 and 2010).
- (e)** Repeat **(d)** using LDA.
- (f)** Repeat **(d)** using QDA.
- (g)** Repeat **(d)** using KNN with $K = 1$.
- (h)** Which of these methods appears to provide the best results on this data?
- (i)** Experiment with different combinations of predictors, including possible transformations and interactions, for each of the methods. Report the variables, method, and associated confusion matrix that appears to provide the best results on the held out data. Note that you should also experiment with values for K in the KNN classifier.

Answer

We will need the data in a variable for ease of use.

```
weeklyDat<-ISLR::Weekly
```

a) Summary Statistics

Text

Most of this segment relies heavily on usage of `dplyr` and especially the `%>%` or pipe operator for readability. The use of the `skimr` package might added more descriptive statistics, but is not covered here.

Basic Summaries

```
weeklyDat %>% str
```

```
# 'data.frame': 1089 obs. of  9 variables:
#  $ Year      : num  1990 1990 1990 1990 1990 1990 1990 1990 1990 1990 ...
#  $ Lag1      : num   0.816 -0.27 -2.576 3.514 0.712 ...
#  $ Lag2      : num   1.572 0.816 -0.27 -2.576 3.514 ...
#  $ Lag3      : num  -3.936 1.572 0.816 -0.27 -2.576 ...
#  $ Lag4      : num  -0.229 -3.936 1.572 0.816 -0.27 ...
#  $ Lag5      : num  -3.484 -0.229 -3.936 1.572 0.816 ...
#  $ Volume     : num   0.155 0.149 0.16 0.162 0.154 ...
#  $ Today      : num  -0.27 -2.576 3.514 0.712 1.178 ...
#  $ Direction: Factor w/ 2 levels "Down","Up": 1 1 2 2 2 1 2 2 2 1 ...
```

We see that there is only one `Factor` , which makes sense.

```
weeklyDat %>% summary
```

```
#      Year      Lag1      Lag2      Lag3
#  Min.   :1990   Min.   :-18.1950   Min.   :-18.1950   Min.   :-18.1950
#  1st Qu.:1995   1st Qu.: -1.1540   1st Qu.: -1.1540   1st Qu.: -1.1580
#  Median :2000   Median :  0.2410   Median :  0.2410   Median :  0.2410
#  Mean   :2000   Mean   :  0.1506   Mean   :  0.1511   Mean   :  0.1472
#  3rd Qu.:2005   3rd Qu.:  1.4050   3rd Qu.:  1.4090   3rd Qu.:  1.4090
#  Max.   :2010   Max.   : 12.0260   Max.   : 12.0260   Max.   : 12.0260
#      Lag4      Lag5      Volume      Today
#  Min.   :-18.1950   Min.   :-18.1950   Min.   :0.08747   Min.   :-18.1950
#  1st Qu.: -1.1580   1st Qu.: -1.1660   1st Qu.:0.33202   1st Qu.: -1.1540
#  Median :  0.2380   Median :  0.2340   Median :1.00268   Median :  0.2410
#  Mean   :  0.1458   Mean   :  0.1399   Mean   :1.57462   Mean   :  0.1499
#  3rd Qu.:  1.4090   3rd Qu.:  1.4050   3rd Qu.:2.05373   3rd Qu.:  1.4050
#  Max.   : 12.0260   Max.   : 12.0260   Max.   :9.32821   Max.   : 12.0260
#  Direction
#  Down:484
#  Up  :605
#
#
#
```

Unique Values

We might also want to know how many unique values are there in each column.

```
weeklyDat %>% apply(unique) %>% apply(length)
```

```
#      Year      Lag1      Lag2      Lag3      Lag4      Lag5      Volume      Today
#      21      1004      1005      1005      1005      1005      1089      1003
# Direction
#      2
```

We note that year has disproportionately lower values, something to keep in mind while constructing models later.

Range

The range of each variable might be useful as well, but we have to ignore the factor.

```
weeklyDat %>% subset(select=-c(Direction)) %>% apply(range)
```

```
#      Year      Lag1      Lag2      Lag3      Lag4      Lag5      Volume      Today
# [1,] 1990 -18.195 -18.195 -18.195 -18.195 -18.195 0.087465 -18.195
# [2,] 2010 12.026 12.026 12.026 12.026 12.026 9.328214 12.026
```

The most interesting thing about this is probably that the `Lag` variables all have the same range, also something to be kept in mind while applying transformations to the variable (if at all).

Mean and Std. Dev

By now we might have a pretty good idea of how this will look, but it is still worth seeing.

```
weeklyDat %>% subset(select=-c(Direction)) %>% apply(mean)
```

```
#      Year      Lag1      Lag2      Lag3      Lag4      Lag5
# 2000.0486685 0.1505849 0.1510790 0.1472048 0.1458182 0.1398926
#      Volume      Today
# 1.5746176 0.1498990
```

As expected, the `Lag` values have almost the same mean, what is a bit interesting though, is that the `Today` variable has roughly the same mean as the `Lag` variables.

```
weeklyDat %>% subset(select=-c(Direction)) %>% apply(sd)
```

```
#      Year      Lag1      Lag2      Lag3      Lag4      Lag5      Volume      Today
# 6.033182 2.357013 2.357254 2.360502 2.360279 2.361285 1.686636 2.356927
```

This is largely redundant in terms of new information.

Correlations

```
weeklyDat %>% subset(select=-c(Direction)) %>% cor
```

```
#      Year      Lag1      Lag2      Lag3      Lag4
# Year      1.00000000 -0.032289274 -0.03339001 -0.03000649 -0.031127923
# Lag1     -0.03228927  1.000000000 -0.07485305  0.05863568 -0.071273876
# Lag2     -0.03339001 -0.074853051  1.00000000 -0.07572091  0.058381535
# Lag3     -0.03000649  0.058635682 -0.07572091  1.00000000 -0.075395865
# Lag4     -0.03112792 -0.071273876  0.05838153 -0.07539587  1.000000000
# Lag5     -0.03051910 -0.008183096 -0.07249948  0.06065717 -0.075675027
# Volume    0.84194162 -0.064951313 -0.08551314 -0.06928771 -0.061074617
# Today     -0.03245989 -0.075031842  0.05916672 -0.07124364 -0.007825873
#      Lag5      Volume      Today
# Year     -0.030519101  0.84194162 -0.032459894
# Lag1     -0.008183096 -0.06495131 -0.075031842
# Lag2     -0.072499482 -0.08551314  0.059166717
# Lag3      0.060657175 -0.06928771 -0.071243639
# Lag4     -0.075675027 -0.06107462 -0.007825873
# Lag5      1.000000000 -0.05851741  0.011012698
# Volume   -0.058517414  1.00000000 -0.033077783
# Today     0.011012698 -0.03307778  1.000000000
```

Useful though this is, it is kind of difficult to work with, in this form, so we might as well programmatic-ally remove strongly correlated data instead.

```
# Uses caret
corrCols=weeklyDat %>% subset(select=-c(Direction)) %>% cor %>% findCorrelation(cutoff=0.8)
reducedDat<-weeklyDat[-c(corrCols)]
reducedDat %>% summary
```

```
#      Year      Lag1      Lag2      Lag3
# Min.   :1990   Min.   :-18.1950   Min.   :-18.1950   Min.   :-18.1950
# 1st Qu.:1995   1st Qu.: -1.1540   1st Qu.: -1.1540   1st Qu.: -1.1580
# Median :2000   Median :  0.2410   Median :  0.2410   Median :  0.2410
# Mean   :2000   Mean   :  0.1506   Mean   :  0.1511   Mean   :  0.1472
# 3rd Qu.:2005   3rd Qu.:  1.4050   3rd Qu.:  1.4090   3rd Qu.:  1.4090
# Max.   :2010   Max.   : 12.0260   Max.   : 12.0260   Max.   : 12.0260
#      Lag4      Lag5      Today      Direction
# Min.   :-18.1950   Min.   :-18.1950   Min.   :-18.1950   Down:484
# 1st Qu.: -1.1580   1st Qu.: -1.1660   1st Qu.: -1.1540   Up :605
# Median :  0.2380   Median :  0.2340   Median :  0.2410
# Mean   :  0.1458   Mean   :  0.1399   Mean   :  0.1499
# 3rd Qu.:  1.4090   3rd Qu.:  1.4050   3rd Qu.:  1.4050
# Max.   : 12.0260   Max.   : 12.0260   Max.   : 12.0260
```

We can see that the `Volume` variable has been dropped, since it evidently is strongly correlated with `Year`. This may or may not be a useful insight, but it is good to keep in mind.

Visualization

We will be using the `ggplot2` library throughout for this segment.

Lets start with some scatter plots in a one v/s all scheme, similar to the methodology [described here](#).

```
weeklyDat %>% subset(select=-c(Direction)) %>% gather(-Year,key="Variable", value="Value") %>%
ggplot(aes(x=Value,y=Year)) +
  geom_point() +
  facet_wrap(~Variable) +
  coord_flip()
```




Figure 1: One v/s all for Direction

That didn't really tell us much which we didn't already get from the `cor()` function, but we can go the whole hog and do this for every variable since we don't have that many in the first place..

```
weeklyDat %>% subset(select=-c(Direction)) %>% pairs
```




Figure 2: Pairs

This is not especially useful, and it is doubtful if more scatter-plots will help at all, so let's move on to box plots.

```
weeklyDat %>% pivot_longer(-c(Direction,Volume,Today,Year),names_to="Lag",values_to="Value") %>%
ggplot(aes(x=Direction,y=Value,fill=Lag)) +
  geom_boxplot()
```




Figure 3: Box plots for Direction

```
weeklyDat %>% pivot_longer(-c(Direction,Volume,Today,Year),names_to="Lag",values_to="Value") %>%
  ggplot(aes(x=Today,y=Value,fill=Lag)) +
  geom_boxplot()
```




Figure 4: More box plots

```
weeklyDat %>% pivot_longer(-c(Direction,Volume,Today,Year),names_to="Lag",values_to="Value") %>%
  ggplot(aes(x=Lag,y=Value,fill=Direction)) +
  geom_boxplot()
```





Figure 5: Lag v/s all

This does summarize our text analysis quite well. Importantly, it tells us that the **Today** value is largely unrelated to the **4 Lag** variables.

A really good-looking box-plot is easy to get with the **caret** library:

```
weeklyDat %>% subset(select=-c(Direction)) %>% featurePlot(
  y = weeklyDat$Direction,
  plot = "box",
  # Pass in options to bwplot()
  scales = list(y = list(relation="free"),
                x = list(rot = 90)),
  auto.key = list(columns = 2))
```




Figure 6: Plots with 'caret'

We might want to [visualize](#) our correlation matrix as well.

```
library(reshape2)
```

```
#  
# Attaching package: 'reshape2'
```

```
# The following object is masked from 'package:tidyr':  
#  
# smiths
```

```
weeklyDat %>% subset(select=-c(Direction)) %>% cor %>% melt %>% ggplot(aes(x=Var1,y=Var2,fill=value)) +  
  geom_tile()
```




Figure 7: Heatmap of the correlation matrix

b) Logistic Regression - Predictor Significance

Lets start with the native `glm` function.

```
glm.fit=glm(Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume, data=weeklyDat, family=binomial)
summary(glm.fit)
```

```
#
# Call:
# glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
#       Volume, family = binomial, data = weeklyDat)
#
# Deviance Residuals:
#      Min       1Q   Median       3Q      Max
# -1.6949  -1.2565   0.9913   1.0849   1.4579
#
# Coefficients:
#              Estimate Std. Error z value Pr(>|z|)
# (Intercept)  0.26686    0.08593   3.106  0.0019 **
# Lag1        -0.04127    0.02641  -1.563  0.1181
# Lag2         0.05844    0.02686   2.175  0.0296 *
# Lag3        -0.01606    0.02666  -0.602  0.5469
# Lag4        -0.02779    0.02646  -1.050  0.2937
# Lag5        -0.01447    0.02638  -0.549  0.5833
# Volume       -0.02274    0.03690  -0.616  0.5377
# ---
# Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#
# (Dispersion parameter for binomial family taken to be 1)
#
#      Null deviance: 1496.2  on 1088  degrees of freedom
# Residual deviance: 1486.4  on 1082  degrees of freedom
# AIC: 1500.4
#
# Number of Fisher Scoring iterations: 4
```

Evidently, only the `Lag2` value is of statistical significance.

It is always of importance to figure out what numeric values R will assign to our factors, and it is best not to guess.

```
contrasts(weeklyDat$Direction)
```

```
#      Up
# Down  0
# Up    1
```

c) Confusion Matrix and Metrics

Essentially:

- Predict the response
- Create an output length vector
- Apply thresholding to obtain labels

```
glm.probs = predict(glm.fit, type = "response")
glm.pred = rep("Up",length(glm.probs))
glm.pred[glm.probs<0.5]="Down"
glm.pred=factor(glm.pred)
confusionMatrix(glm.pred,weeklyDat$Direction)
```

```

# Confusion Matrix and Statistics
#
#           Reference
# Prediction Down  Up
#           Down   54  48
#           Up    430 557
#
#           Accuracy : 0.5611
#           95% CI : (0.531, 0.5908)
#           No Information Rate : 0.5556
#           P-Value [Acc > NIR] : 0.369
#
#           Kappa : 0.035
#
# Mcnemar's Test P-Value : <2e-16
#
#           Sensitivity : 0.11157
#           Specificity : 0.92066
#           Pos Pred Value : 0.52941
#           Neg Pred Value : 0.56434
#           Prevalence : 0.44444
#           Detection Rate : 0.04959
#           Detection Prevalence : 0.09366
#           Balanced Accuracy : 0.51612
#
#           'Positive' Class : Down
#

```

- We have used the `confusionMatrix` function from `caret` ([documented here](#)) instead of displaying the results with `table` and then calculating precision, recall and the rest by hand.

d) Train Test Splits

Although we could have used the indices and passed it to `glm` as the `subset` attribute, it is cleaner to just make subsets instead.

```

weeklyVal<-weeklyDat %>% filter(Year>=2009)
weeklyTrain<-weeklyDat %>% filter(Year<2009)

```

Now we can train a model on our training data.

```

glm.fit=glm(Direction~Lag2,data=weeklyTrain,family=binomial)
summary(glm.fit)

```

```
#
# Call:
# glm(formula = Direction ~ Lag2, family = binomial, data = weeklyTrain)
#
# Deviance Residuals:
#      Min       1Q   Median       3Q      Max
# -1.536  -1.264   1.021   1.091   1.368
#
# Coefficients:
#              Estimate Std. Error z value Pr(>|z|)
# (Intercept)  0.20326    0.06428   3.162  0.00157 **
# Lag2         0.05810    0.02870   2.024  0.04298 *
# ---
# Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#
# (Dispersion parameter for binomial family taken to be 1)
#
#      Null deviance: 1354.7  on 984  degrees of freedom
# Residual deviance: 1350.5  on 983  degrees of freedom
# AIC: 1354.5
#
# Number of Fisher Scoring iterations: 4
```

Having fit our model, we will test the predictions on our held out data.

```
glm.probs = predict(glm.fit,weeklyVal, type = "response")
glm.pred = rep("Up",length(glm.probs))
glm.pred[glm.probs<0.5]="Down"
glm.pred=factor(glm.pred)
confusionMatrix(glm.pred,weeklyVal$Direction)
```

```
# Confusion Matrix and Statistics
#
#              Reference
# Prediction Down Up
#      Down    9  5
#      Up     34 56
#
#              Accuracy : 0.625
#              95% CI : (0.5247, 0.718)
#      No Information Rate : 0.5865
#      P-Value [Acc > NIR] : 0.2439
#
#              Kappa : 0.1414
#
# Mcnemar's Test P-Value : 7.34e-06
#
#              Sensitivity : 0.20930
#              Specificity : 0.91803
#      Pos Pred Value : 0.64286
#      Neg Pred Value : 0.62222
#              Prevalence : 0.41346
#      Detection Rate : 0.08654
#      Detection Prevalence : 0.13462
#      Balanced Accuracy : 0.56367
#
#      'Positive' Class : Down
#
```

We really aren't doing very well with this single variable model as is evident.

e) LDA models

At this stage we could use `MASS` to get the `lda` function, but it would be better to just switch to using `caret`. Note that the `caret` prediction is a label by default, so thresholding needs to be specified differently if required.

```
lda.fit=train(Direction~Lag2,data=weeklyTrain,method="lda")
summary(lda.fit)
```

```
#           Length Class      Mode
# prior      2      -none-    numeric
# counts     2      -none-    numeric
# means      2      -none-    numeric
# scaling    1      -none-    numeric
# lev        2      -none-    character
# svd         1      -none-    numeric
# N           1      -none-    numeric
# call        3      -none-    call
# xNames      1      -none-    character
# problemType 1      -none-    character
# tuneValue   1      data.frame list
# obsLevels   2      -none-    character
# param       0      -none-    list
```

```
predict(lda.fit,weeklyVal) %>% confusionMatrix(weeklyVal$Direction)
```

```
# Confusion Matrix and Statistics
#
#           Reference
# Prediction Down Up
#      Down    9   5
#      Up     34  56
#
#               Accuracy : 0.625
#               95% CI : (0.5247, 0.718)
#      No Information Rate : 0.5865
#      P-Value [Acc > NIR] : 0.2439
#
#               Kappa : 0.1414
#
# Mcnemar's Test P-Value : 7.34e-06
#
#               Sensitivity : 0.20930
#               Specificity : 0.91803
#      Pos Pred Value : 0.64286
#      Neg Pred Value : 0.62222
#               Prevalence : 0.41346
#      Detection Rate : 0.08654
#      Detection Prevalence : 0.13462
#      Balanced Accuracy : 0.56367
#
#      'Positive' Class : Down
#
```

f) QDA models

```
qda.fit=train(Direction~Lag2,data=weeklyTrain,method="qda")
summary(qda.fit)
```

```
#           Length Class      Mode
# prior      2      -none-    numeric
# counts     2      -none-    numeric
# means      2      -none-    numeric
# scaling    2      -none-    numeric
# ldet       2      -none-    numeric
# lev        2      -none-    character
# N           1      -none-    numeric
# call       3      -none-    call
# xNames      1      -none-    character
# problemType 1      -none-    character
# tuneValue   1      data.frame list
# obsLevels   2      -none-    character
# param       0      -none-    list
```

```
predict(qda.fit,weeklyVal) %>% confusionMatrix(weeklyVal$Direction)
```

```
# Confusion Matrix and Statistics
#
#           Reference
# Prediction Down Up
#      Down    0  0
#      Up     43 61
#
#           Accuracy : 0.5865
#           95% CI : (0.4858, 0.6823)
#      No Information Rate : 0.5865
#      P-Value [Acc > NIR] : 0.5419
#
#           Kappa : 0
#
#  Mcnemar's Test P-Value : 1.504e-10
#
#           Sensitivity : 0.0000
#           Specificity : 1.0000
#      Pos Pred Value :      NaN
#      Neg Pred Value : 0.5865
#           Prevalence : 0.4135
#      Detection Rate : 0.0000
#      Detection Prevalence : 0.0000
#      Balanced Accuracy : 0.5000
#
#           'Positive' Class : Down
#
```

This is quite possibly the worst of the lot. As is evident, the model just predicts **Up** no matter what.

g) KNN

`caret` tends to over-zealously retrain models and find the best possible parameters. In this case that is annoying and redundant so we will use the `class` library. We **should really scale our data** before using KNN though.

```
library(class)
set.seed(1)
knn.pred=knn(as.matrix(weeklyTrain$Lag2),as.matrix(weeklyVal$Lag2),weeklyTrain$Direction,k=1)
confusionMatrix(knn.pred,weeklyVal$Direction)
```



```
# Confusion Matrix and Statistics
#
#           Reference
# Prediction Down Up
#      Down   21 30
#      Up    22 31
#
#           Accuracy : 0.5
#           95% CI : (0.4003, 0.5997)
#      No Information Rate : 0.5865
#      P-Value [Acc > NIR] : 0.9700
#
#           Kappa : -0.0033
#
# Mcnemar's Test P-Value : 0.3317
#
#           Sensitivity : 0.4884
#           Specificity : 0.5082
#      Pos Pred Value : 0.4118
#      Neg Pred Value : 0.5849
#           Prevalence : 0.4135
#      Detection Rate : 0.2019
#      Detection Prevalence : 0.4904
#      Balanced Accuracy : 0.4983
#
#      'Positive' Class : Down
#
```

Clearly this model is not doing very well.

h) Model Selection

We will first get the ROC curves.

```
library(pROC)
```

```
# Type 'citation("pROC")' for a citation.
```

```
#
# Attaching package: 'pROC'
```

```
# The following objects are masked from 'package:stats':
#
#      cov, smooth, var
```

```
knnROC<-
roc(predictor=as.numeric(knn.pred),response=weeklyVal$Direction,levels=rev(levels(weeklyVal$Direction)))
```

```
# Setting direction: controls < cases
```

```
logiROC<-roc(predictor=as.numeric(predict(glm.fit,weeklyVal)),response=weeklyVal$Direction)
```

```
# Setting levels: control = Down, case = Up
```

```
# Setting direction: controls > cases
```

```
ldaROC<-roc(predictor=as.numeric(predict(lda.fit,weeklyVal)),response=weeklyVal$Direction)
```

```
# Setting levels: control = Down, case = Up
```

```
# Setting direction: controls < cases
```

```
qdaROC<-roc(predictor=as.numeric(predict(qda.fit,weeklyVal)),response=weeklyVal$Direction)
```

```
# Setting levels: control = Down, case = Up
```

```
# Setting direction: controls < cases
```

Now to plot them.

```
ggroc(list(KNN=knnROC,Logistic=logiROC,LDA=ldaROC,QDA=qdaROC))
```




Figure 8: ROC curves for Weekly data

To compare models with `caret` it is easy to refit the logistic and knn models in the caret formulation.

```
knnCaret=train(Direction~Lag2,data=weeklyTrain,method="knn")
```

However, the KNN model is the best parameter model.

```
resmod <- resamples(list(lda=lda.fit, qda=qda.fit, KNN=knnCaret))
summary(resmod)
```

```
#
# Call:
# summary.resamples(object = resmod)
#
# Models: lda, qda, KNN
# Number of resamples: 25
#
# Accuracy
#           Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
# lda 0.5043228 0.5344353 0.5529101 0.5500861 0.5683060 0.5846995    0
# qda 0.5044248 0.5204360 0.5307263 0.5326785 0.5462428 0.5777778    0
# KNN 0.4472222 0.5082873 0.5240642 0.5168327 0.5302198 0.5485714    0
#
# Kappa
#           Min.      1st Qu.      Median      Mean      3rd Qu.      Max.
# lda -0.02618939 -0.003638168 0.005796908 0.007801904 0.01635328 0.05431238
# qda -0.06383592 -0.005606123 0.000000000 -0.003229697 0.000000000 0.03606344
# KNN -0.11297539 0.004168597 0.024774647 0.016171229 0.04456142 0.07724439
#           NA's
# lda      0
# qda      0
# KNN      0
```

```
bwplot(resmod)
```





Figure 9: Caret plots for comparison

```
dotplot(resmod)
```



Kappa or Cohen's Kappa is essentially classification accuracy, normalized at the baseline of random chance. It is a more useful measure to use on problems that have imbalanced classes. There's more on model selection [here](#).

i) Further Tuning

Do note the `caret` defaults.

```
fitControl <- trainControl(# 10-fold CV
  method = "repeatedcv",
  number = 10,
  # repeated ten times
  repeats = 10)
```

Logistic

```
glm2.fit=glm(Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume, data=weeklyDat, family=binomial)

glm2.probs = predict(glm2.fit,weeklyVal, type = "response")
glm2.pred = rep("Up",length(glm2.probs))
glm2.pred[glm2.probs<0.5]="Down"
glm2.pred=factor(glm2.pred)
confusionMatrix(glm2.pred,weeklyVal$Direction)
```

```
# Confusion Matrix and Statistics
#
#               Reference
# Prediction Down Up
#      Down   17  13
#      Up    26  48
#
#               Accuracy : 0.625
#               95% CI : (0.5247, 0.718)
#      No Information Rate : 0.5865
#      P-Value [Acc > NIR] : 0.24395
#
#               Kappa : 0.1907
#
#  Mcnemar's Test P-Value : 0.05466
#
#               Sensitivity : 0.3953
#               Specificity : 0.7869
#      Pos Pred Value : 0.5667
#      Neg Pred Value : 0.6486
#      Prevalence : 0.4135
#      Detection Rate : 0.1635
#      Detection Prevalence : 0.2885
#      Balanced Accuracy : 0.5911
#
#      'Positive' Class : Down
#
```

QDA

```
qdaCaret=train(Direction~Lag2+Lag4,data=weeklyTrain,method="qda",trainControl=fitControl)
```

```
summary(qdaCaret)
```

```
#               Length Class      Mode
# prior          2      -none-   numeric
# counts         2      -none-   numeric
# means          4      -none-   numeric
# scaling        8      -none-   numeric
# ldet           2      -none-   numeric
# lev            2      -none-   character
# N              1      -none-   numeric
# call           4      -none-   call
# xNames         2      -none-   character
# problemType    1      -none-   character
# tuneValue      1      data.frame list
# obsLevels      2      -none-   character
# param          1      -none-   list
```

```
predict(qdaCaret,weeklyVal) %>% confusionMatrix(weeklyVal$Direction)
```

```
# Confusion Matrix and Statistics
#
#           Reference
# Prediction Down Up
#      Down    9 14
#      Up     34 47
#
#           Accuracy : 0.5385
#           95% CI : (0.438, 0.6367)
#      No Information Rate : 0.5865
#      P-Value [Acc > NIR] : 0.863079
#
#           Kappa : -0.0217
#
#  Mcnemar's Test P-Value : 0.006099
#
#      Sensitivity : 0.20930
#      Specificity : 0.77049
#      Pos Pred Value : 0.39130
#      Neg Pred Value : 0.58025
#      Prevalence : 0.41346
#      Detection Rate : 0.08654
#      Detection Prevalence : 0.22115
#      Balanced Accuracy : 0.48990
#
#      'Positive' Class : Down
#
```

LDA

```
ldaCaret=train(Direction~Lag2+Lag1+Year,data=weeklyTrain,method="lda",trainControl=fitControl)
```

```
summary(ldaCaret)
```

```
#           Length Class      Mode
# prior         2    -none-   numeric
# counts        2    -none-   numeric
# means         6    -none-   numeric
# scaling       3    -none-   numeric
# lev           2    -none-   character
# svd            1    -none-   numeric
# N             1    -none-   numeric
# call          4    -none-   call
# xNames        3    -none-   character
# problemType  1    -none-   character
# tuneValue     1    data.frame list
# obsLevels     2    -none-   character
# param         1    -none-   list
```

```
predict(ldaCaret,weeklyVal) %>% confusionMatrix(weeklyVal$Direction)
```

```

# Confusion Matrix and Statistics
#
#           Reference
# Prediction Down Up
#      Down   20  19
#      Up    23  42
#
#           Accuracy : 0.5962
#           95% CI : (0.4954, 0.6913)
#      No Information Rate : 0.5865
#      P-Value [Acc > NIR] : 0.4626
#
#           Kappa : 0.1558
#
# Mcnemar's Test P-Value : 0.6434
#
#           Sensitivity : 0.4651
#           Specificity : 0.6885
#      Pos Pred Value : 0.5128
#      Neg Pred Value : 0.6462
#           Prevalence : 0.4135
#      Detection Rate : 0.1923
#      Detection Prevalence : 0.3750
#      Balanced Accuracy : 0.5768
#
#      'Positive' Class : Down
#

```

KNN

Honestly, again, this should be scaled. Plot `knn` with the best parameters.

```
plot(knnCaret)
```





Figure 10: KNN statistics

Evidently, the accuracy increases with an increase in the number of neighbors considered.

```
plot(knnCaret, print.thres = 0.5, type="S")
```




Figure 11: Visualizing thresholds for KNN

However this shows that we don't actually get much of an increase in accuracy anyway.

Question 4.11 - Pages 171-172

In this problem, you will develop a model to predict whether a given car gets high or low gas mileage based on the Auto data set.

(a) Create a binary variable, `mpg01`, that contains a 1 if `mpg` contains a value above its median, and a 0 if `mpg` contains a value below its median. You can compute the median using the `median()` function. Note you may find it helpful to use the `data.frame()` function to create a single data set containing both `mpg01` and the other Auto variables.

(b) Explore the data graphically in order to investigate the association between `mpg01` and the other features. Which of the other features seem most likely to be useful in predicting `mpg01`? Scatter-plots and boxplots may be useful tools to answer this question. Describe your findings.

(c) Split the data into a training set and a test set.

(d) Perform LDA on the training data in order to predict `mpg01` using the variables that seemed most associated with `mpg01` in **(b)**. What is the test error of the model obtained?

(e) Perform QDA on the training data in order to predict `mpg01` using the variables that seemed most associated with `mpg01` in **(b)**. What is the test error of the model obtained?

(f) Perform logistic regression on the training data in order to predict `mpg01` using the variables that seemed most associated with `mpg01` in **(b)**. What is the test error of the model obtained?

(g) Perform KNN on the training data, with several values of K , in order to predict `mpg01`. Use only the variables that seemed most associated with `mpg01` in **(b)**. What test errors do you obtain? Which value of K seems to perform the best on this data set?

Answer

```
autoDat<- ISLR::Auto
```

a) Binary Variable

```
autoDat$mpg %>% sort() %>% median()
```

```
# [1] 22.75
```

Now we can get a new variable from that.

```
newDat=autoDat
newDat$mpg01 <- ifelse(autoDat$mpg<autoDat$mpg %>% sort() %>% median(),0,1) %>% factor()
```

Note that the `ifelse` command takes a truthy function, value when false, value when true, but does not return a factor automatically so we piped it to `factor` to ensure it is factorial.

b) Visual Exploration

Some box-plots:

```
newDat %>% pivot_longer(-c(mpg01,name),names_to="Params",values_to="Value") %>% ggplot(aes(x=mpg01,y=Value)) +
  geom_boxplot() +
  facet_wrap(~ Params, scales = "free_y")
```




Figure 12: Box plots

With some scatter plots as well:

```
newDat %>% pivot_longer(-c(mpg01,name,weight),names_to="Params",values_to="Value") %>%
ggplot(aes(x=weight,y=Value,color=mpg01)) +
  geom_point() +
  facet_wrap(~ Params, scales = "free_y")
```





Figure 13: Scatter plots

Clearly, `origin`, `year` and `cylinder` are essentially not very relevant numerically for the regression lines and confidence intervals.

```
newDat %>% select(-year,-origin,-cylinders) %>% pivot_longer(-
c(mpg01,name,mpg),names_to="Params",values_to="Value") %>% ggplot(aes(x=mpg,y=Value,color=mpg01)) +
  geom_point() +
  geom_smooth(method=lm) +
  facet_wrap(~ Params, scales = "free_y")
```



c) Train-Test Split

We can split our data [very easily](#) with `caret`. It is important to remember that for factors, random sampling occurs within each class to preserve the overall class distribution of the data.

```
set.seed(1984)
trainInd <- createDataPartition(newDat$mpg01, # Factor, so class sampling
                                p=0.7, # 70-30 train-test
                                list=FALSE, # No lists
                                times=1) # No bootstrap

autoTrain<-newDat[trainInd,]
autoTest<-newDat[-trainInd,]
```

d) LDA with Significant Variables

Whenever I see significant I think correlation, so let's take a look at that.

```
newDat %>% select(-mpg01,-name) %>% cor
```

```
#           mpg  cylinders displacement horsepower    weight
# mpg      1.0000000 -0.7776175   -0.8051269 -0.7784268 -0.8322442
# cylinders -0.7776175  1.0000000    0.9508233  0.8429834  0.8975273
# displacement -0.8051269  0.9508233    1.0000000  0.8972570  0.9329944
# horsepower  -0.7784268  0.8429834    0.8972570  1.0000000  0.8645377
# weight      -0.8322442  0.8975273    0.9329944  0.8645377  1.0000000
# acceleration 0.4233285 -0.5046834   -0.5438005 -0.6891955 -0.4168392
# year        0.5805410 -0.3456474   -0.3698552 -0.4163615 -0.3091199
# origin      0.5652088 -0.5689316   -0.6145351 -0.4551715 -0.5850054
#
#           acceleration    year    origin
# mpg      0.4233285  0.5805410  0.5652088
# cylinders -0.5046834 -0.3456474 -0.5689316
# displacement -0.5438005 -0.3698552 -0.6145351
# horsepower  -0.6891955 -0.4163615 -0.4551715
# weight      -0.4168392 -0.3091199 -0.5850054
# acceleration 1.0000000  0.2903161  0.2127458
# year        0.2903161  1.0000000  0.1815277
# origin      0.2127458  0.1815277  1.0000000
```

```
newDat %>% length
```

```
# [1] 10
```

Now lets quickly see what it looks like with correlated values removed.

```
corrCols2=newDat %>% select(-mpg01,-name) %>% cor %>% findCorrelation(cutoff=0.85)
newRed<-newDat[-c(corrCols2)]
newRed %>% summary
```

```
#           mpg           weight      acceleration          year          origin
# Min.      : 9.00   Min.      :1613   Min.      : 8.00   Min.      :70.00   Min.      :1.000
# 1st Qu.:17.00   1st Qu.:2225   1st Qu.:13.78   1st Qu.:73.00   1st Qu.:1.000
# Median :22.75   Median :2804   Median :15.50   Median :76.00   Median :1.000
# Mean      :23.45   Mean      :2978   Mean      :15.54   Mean      :75.98   Mean      :1.577
# 3rd Qu.:29.00   3rd Qu.:3615   3rd Qu.:17.02   3rd Qu.:79.00   3rd Qu.:2.000
# Max.      :46.60   Max.      :5140   Max.      :24.80   Max.      :82.00   Max.      :3.000
#
#
#           name      mpg01
# amc matador      : 5    0:196
# ford pinto       : 5    1:196
# toyota corolla   : 5
# amc gremlin      : 4
# amc hornet       : 4
# chevrolet chevette: 4
# (Other)          :365
```

Inherent in this discussion is the fact that I consider what is correlated to `mpg` to be a good indicator of what will help `mpg01` for obvious reasons.

Now we can just use the columns we found with `findCorrelation`.

```
corrCols2 %>% print
```

```
# [1] 3 4 2
```

```
names(newDat)
```

```
# [1] "mpg"          "cylinders"    "displacement" "horsepower"   "weight"
# [6] "acceleration" "year"         "origin"       "name"         "mpg01"
```

```
autoLDA=train(mpg01~cylinders+displacement+horsepower,data=autoTrain,method="lda")
valScoreLDA=predict(autoLDA,autoTest)
```

Now we can check the statistics.

```
confusionMatrix(valScoreLDA,autoTest$mpg01)
```


```
# Confusion Matrix and Statistics
#
#           Reference
# Prediction  0   1
#           0 56  2
#           1  2 56
#
#               Accuracy : 0.9655
#               95% CI : (0.9141, 0.9905)
#       No Information Rate : 0.5
#       P-Value [Acc > NIR] : <2e-16
#
#               Kappa : 0.931
#
#  Mcnemar's Test P-Value : 1
#
#               Sensitivity : 0.9655
#               Specificity : 0.9655
#               Pos Pred Value : 0.9655
#               Neg Pred Value : 0.9655
#               Prevalence : 0.5000
#               Detection Rate : 0.4828
#       Detection Prevalence : 0.5000
#       Balanced Accuracy : 0.9655
#
#       'Positive' Class : 0
#
```

That is an amazingly accurate model.

```
auto_ldaROC<-roc(predictor=as.numeric(valScoreLDA),response=autoTest$mpg01,levels=levels(autoTest$mpg01))
```

```
# Setting direction: controls < cases
```

```
ggroc(auto_ldaROC)
```

e) QDA with Significant Variables

Same deal as before.

```
autoQDA=train(mpg01~cylinders+displacement+horsepower,data=autoTrain,method="qda")
valScoreQDA=predict(autoQDA,autoTest)
```

Now we can check the statistics.

```
confusionMatrix(valScoreQDA,autoTest$mpg01)
```

```

# Confusion Matrix and Statistics
#
#           Reference
# Prediction  0   1
#           0 56  2
#           1  2 56
#
#           Accuracy : 0.9655
#           95% CI : (0.9141, 0.9905)
#       No Information Rate : 0.5
#       P-Value [Acc > NIR] : <2e-16
#
#           Kappa : 0.931
#
#  Mcnemar's Test P-Value : 1
#
#           Sensitivity : 0.9655
#           Specificity : 0.9655
#       Pos Pred Value : 0.9655
#       Neg Pred Value : 0.9655
#           Prevalence : 0.5000
#       Detection Rate : 0.4828
#   Detection Prevalence : 0.5000
#       Balanced Accuracy : 0.9655
#
#       'Positive' Class : 0
#

```

```

auto_qdaROC<-roc(predictor=as.numeric(valScoreQDA),response=autoTest$mpg01,levels=levels(autoTest$mpg01))

```

```


# Setting direction: controls < cases

```

```

ggroc(auto_qdaROC)

```



OK, this is weird enough to check if it isn't some sort of artifact.

```
autoQDA2=train(mpg01~horsepower, data=autoTrain,method='qda')
valScoreQDA2=predict(autoQDA2, autoTest)
confusionMatrix(valScoreQDA2,autoTest$mpg01)
```

```

# Confusion Matrix and Statistics
#
#           Reference
# Prediction  0   1
#           0  42   3
#           1  16  55
#
#           Accuracy : 0.8362
#           95% CI : (0.7561, 0.8984)
#           No Information Rate : 0.5
#           P-Value [Acc > NIR] : 4.315e-14
#
#           Kappa : 0.6724
#
# Mcnemar's Test P-Value : 0.005905
#
#           Sensitivity : 0.7241
#           Specificity : 0.9483
#           Pos Pred Value : 0.9333
#           Neg Pred Value : 0.7746
#           Prevalence : 0.5000
#           Detection Rate : 0.3621
#           Detection Prevalence : 0.3879
#           Balanced Accuracy : 0.8362
#
#           'Positive' Class : 0
#

```

OK, so the model isn't completely creepily correct all the time. In this case we should probably think about what is going on. I would think it is because of the nature of the `train-test` split we performed. We have ensured during the sampling of our data that the train and test sets contain the SAME distribution (assumed). So that's why our training result and test results are both incredibly good. They're essentially the same thing.

In fact, this is the perfect time to consider a validation set, just to see what the models are really doing. Won't get into it right now though.

f) Logistic with Significant Variables

```
glmAuto.fit=glm(mpg01~cylinders+displacement+horsepower, data=autoTrain, family=binomial)
```

```

glmAuto.probs = predict(glmAuto.fit,autoTest, type = "response")
glmAuto.pred = rep(1,length(glmAuto.probs))
glmAuto.pred[glmAuto.probs<0.5]=0
glmAuto.pred=factor(glmAuto.pred)
confusionMatrix(glmAuto.pred,autoTest$mpg01)

```

```

# Confusion Matrix and Statistics
#
#           Reference
# Prediction  0   1
#           0 56   4
#           1   2 54
#
#           Accuracy : 0.9483
#           95% CI : (0.8908, 0.9808)
#           No Information Rate : 0.5
#           P-Value [Acc > NIR] : <2e-16
#
#           Kappa : 0.8966
#
# Mcnemar's Test P-Value : 0.6831
#
#           Sensitivity : 0.9655
#           Specificity : 0.9310
#           Pos Pred Value : 0.9333
#           Neg Pred Value : 0.9643
#           Prevalence : 0.5000
#           Detection Rate : 0.4828
#           Detection Prevalence : 0.5172
#           Balanced Accuracy : 0.9483
#
#           'Positive' Class : 0
#

```


g) KNN Modeling

Scale the parameters later.

```
knnAuto=train(mpg01~cylinders+displacement+horsepower,data=autoTrain,method="knn")
```


Plot `knn` with the best parameters.

```
plot(knnCaret)
```



Evidently, the accuracy increases with an increase in the number of neighbors considered.

```
plot(knnAuto, print.thres = 0.5, type="S")
```



So we can see that 5 neighbors is a good compromise.

Question 4.12 - Pages 172-173

This problem involves writing functions.

(a) Write a function, `Power()`, that prints out the result of raising 2 to the 3rd power. In other words, your function should compute 2^3 and print out the results.

Hint: Recall that x^a raises x to the power a . Use the `print()` function to output the result.

(b) Create a new function, `Power2()`, that allows you to pass any two numbers, `x` and `a`, and prints out the value of x^a . You can do this by beginning your function with the line

```
Power2=function(x,a){}
```

You should be able to call your function by entering, for instance,

```
Power2(3,8)
```

on the command line. This should output the value of 3^8 , namely, 6,651.

(c) Using the `Power2()` function that you just wrote, compute 10^3 , 8^{17} , and 131^3 .

(d) Now create a new function, `Power3()`, that actually *returns* the result `x^a` as an R object, rather than simply printing it to the screen. That is, if you store the value `x^a` in an object called `result` within your function, then you can simply `return()` this result, using the following line:

```
return(result)
```

The line above should be the last line in your function, before the `}` symbol.

(e) Now using the `Power3()` function, create a plot of $f(x) = x^2$. The x-axis should display a range of integers from 1 to 10, and the y-axis should display x^2 . Label the axes appropriately, and use an appropriate title for the figure. Consider displaying either the x-axis, the y-axis, or both on the log-scale. You can do this by using `log='x'`, `log='y'`, or `log='xy'` as arguments to the `plot()` function.

(f) Create a function, `PlotPower()`, that allows you to create a plot of `x` against `x^a` for a fixed `a` and for a range of values of `x`. For instance, if you call

```
PlotPower (1:10 ,3)
```

then a plot should be created with an x-axis taking on values 1, 2, ..., 10 and a y-axis taking on values $1^3, 2^3, \dots, 10^3$

Answer

a) Create a Squaring Function

```
Power=function(x){print(2^x)}  
Power(3)
```

```
# [1] 8
```

b) Generalizing Power to arbitrary numbers

```
Power2=function(x,a){print(x^a)}
```

```
Power2(3,8)
```

```
# [1] 6561
```

c) Random Testing of Power2

```
Power2(10,3)
```

```
# [1] 1000
```

```
Power2(8,17)
```



```
# [1] 2.2518e+15
```

```
Power2(131,2)
```

```
# [1] 17161
```

d) Return a value

```
Power3=function(x,a){return(x^a)}
```

e) Plot something with Power3


Actually now would be a good place to introduce `LaTeX` labeling.

```
#install.packages("latex2exp")  
library(latex2exp)
```

No log scale.

```
qplot(x=seq(1,10),y=Power3(seq(1,10),2)) + ggtitle("Function without a log scale") +  
  geom_point() + xlab("X") + ylab(TeX("$X^2$"))
```


Function without a log scale



With a log scale.

```
qplot(x=seq(1,10),y=Power3(seq(1,10),2)) + ggtitle("Function with a log scale") +  
  geom_point() + xlab("X") + ylab(TeX("$X^2$")) + scale_y_log10()
```


Function with a log scale



f) PlotPower Function

```
PlotPower=function(xrange,pow){return(qplot(x=xrange,y=Power3(xrange,pow)))}
```

```
plotter<-PlotPower(1:10,3)  
plotter
```



The [R Cookbook](#) is quite neat for some simple tasks like this.

Question 4.13 - Pages 173

Using the `Boston` data set, fit classification models in order to predict whether a given suburb has a crime rate above or below the median. Explore logistic regression, LDA, and KNN models using various subsets of the predictors. Describe your findings.

Answer

OK, to speed this up, I will simply run through all the work done on the `Auto` set. Recall that details about this data-set are [also here](#).

```
boston<-MASS::Boston
```

- Check unique values

```
boston %>% apply(unique) %>% apply(length)
```

```
#   crim    zn   indus   chas   nox    rm    age    dis    rad    tax
#   504     26    76     2     81   446   356   412     9    66
# ptratio  black  lstat   medv
#    46     357   455    229
```

CHAS is of course something which should be a factor, and with **RAD** having only **9** levels, I'm inclined to make it a factor as well.

```
boston<-boston %>% mutate(rad=factor(rad),chas=factor(chas))
```


- Make a median variable

```
boston$highCrime<- ifelse(boston$crim<boston$crim %>% median(),0,1) %>% factor()
```

- Take a look at the data

Some box-plots:

```
boston %>% pivot_longer(-c(rad,chas,highCrime),names_to="Param",values_to="Value") %>%
ggplot(aes(x=highCrime,y=Value,fill=chas)) +
  geom_boxplot()+
  facet_wrap(~Param,scales="free_y")
```



It is surprising, but evidently the **CHAS** variable is strangely relevant. 1 implies the tract bounds the river, otherwise 0.

- Correlations

```
boston %>% select(-c(rad,chas,highCrime)) %>% cor
```

#	crim	zn	indus	nox	rm	age
# crim	1.0000000	-0.2004692	0.4065834	0.4209717	-0.2192467	0.3527343
# zn	-0.2004692	1.0000000	-0.5338282	-0.5166037	0.3119906	-0.5695373
# indus	0.4065834	-0.5338282	1.0000000	0.7636514	-0.3916759	0.6447785
# nox	0.4209717	-0.5166037	0.7636514	1.0000000	-0.3021882	0.7314701
# rm	-0.2192467	0.3119906	-0.3916759	-0.3021882	1.0000000	-0.2402649
# age	0.3527343	-0.5695373	0.6447785	0.7314701	-0.2402649	1.0000000
# dis	-0.3796701	0.6644082	-0.7080270	-0.7692301	0.2052462	-0.7478805
# tax	0.5827643	-0.3145633	0.7207602	0.6680232	-0.2920478	0.5064556
# ptratio	0.2899456	-0.3916785	0.3832476	0.1889327	-0.3555015	0.2615150
# black	-0.3850639	0.1755203	-0.3569765	-0.3800506	0.1280686	-0.2735340
# lstat	0.4556215	-0.4129946	0.6037997	0.5908789	-0.6138083	0.6023385
# medv	-0.3883046	0.3604453	-0.4837252	-0.4273208	0.6953599	-0.3769546
#	dis	tax	ptratio	black	lstat	medv
# crim	-0.3796701	0.5827643	0.2899456	-0.3850639	0.4556215	-0.3883046
# zn	0.6644082	-0.3145633	-0.3916785	0.1755203	-0.4129946	0.3604453
# indus	-0.7080270	0.7207602	0.3832476	-0.3569765	0.6037997	-0.4837252
# nox	-0.7692301	0.6680232	0.1889327	-0.3800506	0.5908789	-0.4273208
# rm	0.2052462	-0.2920478	-0.3555015	0.1280686	-0.6138083	0.6953599
# age	-0.7478805	0.5064556	0.2615150	-0.2735340	0.6023385	-0.3769546
# dis	1.0000000	-0.5344316	-0.2324705	0.2915117	-0.4969958	0.2499287
# tax	-0.5344316	1.0000000	0.4608530	-0.4418080	0.5439934	-0.4685359
# ptratio	-0.2324705	0.4608530	1.0000000	-0.1773833	0.3740443	-0.5077867
# black	0.2915117	-0.4418080	-0.1773833	1.0000000	-0.3660869	0.3334608
# lstat	-0.4969958	0.5439934	0.3740443	-0.3660869	1.0000000	-0.7376627
# medv	0.2499287	-0.4685359	-0.5077867	0.3334608	-0.7376627	1.0000000

Now, unsurprisingly, there's nothing which is really strongly correlated here for some reason.

- Train test splits

```
set.seed(1984)
trainIndCri <- createDataPartition(boston$highCrime, # Factor, so class sampling
                                   p=0.7, # 70-30 train-test
                                   list=FALSE, # No lists
                                   times=1) # No bootstrap

bostonTrain<-boston[trainIndCri,]
bostonTest<-boston[-trainIndCri,]
```

- Make a bunch of models

```
glmBos.fit=glm(highCrime~., data=bostonTrain, family=binomial)
```

```
# Warning: glm.fit: algorithm did not converge
```

```
# Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
glmBos.probs = predict(glmBos.fit,bostonTest, type = "response")
glmBos.pred = rep(1,length(glmBos.probs))
glmBos.pred[glmBos.probs<0.5]=0
glmBos.pred=factor(glmBos.pred)
confusionMatrix(glmBos.pred,bostonTest$highCrime)
```

```

# Confusion Matrix and Statistics
#
#           Reference
# Prediction  0   1
#           0 68   6
#           1   7 69
#
#           Accuracy : 0.9133
#           95% CI : (0.8564, 0.953)
#           No Information Rate : 0.5
#           P-Value [Acc > NIR] : <2e-16
#
#           Kappa : 0.8267
#
# Mcnemar's Test P-Value : 1
#
#           Sensitivity : 0.9067
#           Specificity : 0.9200
#           Pos Pred Value : 0.9189
#           Neg Pred Value : 0.9079
#           Prevalence : 0.5000
#           Detection Rate : 0.4533
#           Detection Prevalence : 0.4933
#           Balanced Accuracy : 0.9133
#
#           'Positive' Class : 0
#

```

```

bostonLDA=train(highCrime~.,data=bostonTrain,method='lda')
bostonQDA=train(highCrime~tax+crim,data=bostonTrain,method='qda')
bostonKNN=train(highCrime~.,data=bostonTrain,preProcess = c("center","scale"),method='knn')

```

```

bLDAp=predict(bostonLDA,bostonTest)
bQDAp=predict(bostonQDA,bostonTest)
bKNNp=predict(bostonKNN,bostonTest)

```

```

confusionMatrix(bLDAp,bostonTest$highCrime)

```



```
# Confusion Matrix and Statistics
#
#           Reference
# Prediction  0   1
#           0 72  6
#           1  3 69
#
#           Accuracy : 0.94
#           95% CI : (0.8892, 0.9722)
#       No Information Rate : 0.5
#       P-Value [Acc > NIR] : <2e-16
#
#           Kappa : 0.88
#
# Mcnemar's Test P-Value : 0.505
#
#           Sensitivity : 0.9600
#           Specificity : 0.9200
#       Pos Pred Value : 0.9231
#       Neg Pred Value : 0.9583
#           Prevalence : 0.5000
#       Detection Rate : 0.4800
#       Detection Prevalence : 0.5200
#       Balanced Accuracy : 0.9400
#
#       'Positive' Class : 0
#
```

```
confusionMatrix(bQDAp,bostonTest$highCrime)
```

```
# Confusion Matrix and Statistics
#
#           Reference
# Prediction  0   1
#           0 73  5
#           1  2 70
#
#           Accuracy : 0.9533
#           95% CI : (0.9062, 0.981)
#       No Information Rate : 0.5
#       P-Value [Acc > NIR] : <2e-16
#
#           Kappa : 0.9067
#
# Mcnemar's Test P-Value : 0.4497
#
#           Sensitivity : 0.9733
#           Specificity : 0.9333
#       Pos Pred Value : 0.9359
#       Neg Pred Value : 0.9722
#           Prevalence : 0.5000
#       Detection Rate : 0.4867
#       Detection Prevalence : 0.5200
#       Balanced Accuracy : 0.9533
#
#       'Positive' Class : 0
#
```

```
confusionMatrix(bKNNp,bostonTest$highCrime)
```

```


# Confusion Matrix and Statistics
#
#           Reference
# Prediction  0   1
#           0 74   6
#           1   1 69
#
#           Accuracy : 0.9533
#           95% CI : (0.9062, 0.981)
#           No Information Rate : 0.5
#           P-Value [Acc > NIR] : <2e-16
#
#           Kappa : 0.9067
#
# Mcnemar's Test P-Value : 0.1306
#
#           Sensitivity : 0.9867
#           Specificity : 0.9200
#           Pos Pred Value : 0.9250
#           Neg Pred Value : 0.9857
#           Prevalence : 0.5000
#           Detection Rate : 0.4933
#           Detection Prevalence : 0.5333
#           Balanced Accuracy : 0.9533
#
#           'Positive' Class : 0
#

```


Clearly in this particular case, an LDA model seems to be working out the best for this data when trained on all the parameters, though Logistic Regression is doing quite well too.

- Notes on KNN

```
plot(bostonKNN)
```



```
plot(bostonKNN, print.thres = 0.5, type="S")
```



- Comparison

Finally, we will quickly plot some indicative measures.

```
knnBosROC<- roc(predictor=as.numeric(bKNNp), response=bostonTest$highCrime)
```

```
# Setting levels: control = 0, case = 1
```

```
# Setting direction: controls < cases
```

```
logiBosROC<- roc(predictor=as.numeric(glmBos.probs), response=bostonTest$highCrime)
```

```
# Setting levels: control = 0, case = 1  
# Setting direction: controls < cases
```

```
ldaBosROC<- roc(predictor=as.numeric(bLDAP), response=bostonTest$highCrime)
```

```
# Setting levels: control = 0, case = 1  
# Setting direction: controls < cases
```

```
qdaBosROC<- roc(predictor=as.numeric(bQDAp), response=bostonTest$highCrime)
```

```
# Setting levels: control = 0, case = 1  
# Setting direction: controls < cases
```

```
ggroc(list(KNN=knnBosROC, Logistic=logiBosROC, LDA=ldaBosROC, QDA=qdaBosROC))
```




Figure 14: plot of chunk unnamed-chunk-87

OK, one of the reasons why these models do so well is because they are all assuming an equal distribution of train and test classes, and they use `crim` itself as a predictor. This is no doubt a strong reason why these models uniformly perform so well. I'd say 5 is the best option.

