

Trivial LLMs on old hardware

Rohit Goswami · 2025-11-22 · ai · hardware · workflow

Apropos nothing. R1 represents the remote machine. L1 denotes the local machine. Tailscale connects the two.

```
# R1, assuming Ollama is installed
pixi global install --environment llmlocal ollama
# or with a GPU
pixi global install --environment llmlocal "ollama *cuda*"
export OLLAMA_HOST="0.0.0.0"
ollama serve
# another tab
ollama run MODEL_NAME
```

A preliminary check on the local machine:

```
curl http://localhost:11434/api/tags
```

Followed by some more reasonable usage:

```
# L1
ssh -N -L 11434:127.0.0.1:11434 $R1_HOST
export OLLAMA_API_BASE=http://localhost:11434
pixi global install --environment llmlocal aider-chat gpustat
aider --model ollama/qwen3-coder:480b-cloud blah.py
# different tab
gpustat -i 1
```

Tunneling the API port over SSH via Tailscale ensures a secure connection without exposing the Ollama instance to the open web. This workflow facilitates the usage of substantial models on lightweight clients, effectively offloading the heavy lifting to dedicated hardware, and making good use of older machines.

Also see this post for [more details](#).